



# 强化学习原理及应用 Reinforcement Learning (RL): Theories & Applications

*DCS6289 Spring 2022*

Yucong Zhang (张宇聪)

School of Computer Science and Engineering  
Sun Yat-Sen University

# Lecture 10: Multi-Agent RL

17<sup>th</sup> May. 2022

## ❑ Independent Learning

- ❑ IQL
- ❑ IL with parameter sharing

## ❑ Learning cooperation

- ❑ MADDPG
- ❑ COMA
- ❑ VDN
- ❑ QMIX

# Multi-Agent RL

## □ Independent Learning

□ Evaluate single-agent DRL algorithm in multiagent settings.

## □ IQL

- ✓ Study emergent cooperative and competitive strategies between multiple agents controlled by autonomous deep Q-Networks.
- ✓ Use Atari video games as the environment.
- ✓ Explore how two agents behave and interact in complex environment when trained with different rewarding schemes.



### Rewarding Schemes

1. Fully Competitive
2. Fully Cooperative
3. Transition between Cooperation and Competition

## □ IQL

- Score More than the Opponent (**Fully Competitive**)
  - The player who last touches the outgoing ball gets a plus point, and the player losing the ball a minus point.

	Left player scores	Right player scores
Left player reward	+1	-1
Right player reward	-1	+1

- Loosing the Ball Penalizes Both Players (**Fully Cooperative**)
  - Agents need to learn to keep the ball in the game for as long as possible.
  - Penalizing both of the players whenever the ball goes out of play.

	Left player scores	Right player scores
Left player reward	-1	-1
Right player reward	-1	-1

# Multi-Agent RL

## □ IQL

- Transition Between Cooperation and Competition
  - The fully competitive and fully cooperative cases both penalize loosing the ball equally. What differentiates the two strategies are the values on the main diagonal of the reward matrix.
  - Allow this reward value  $\rho$  to change gradually from -1 to +1

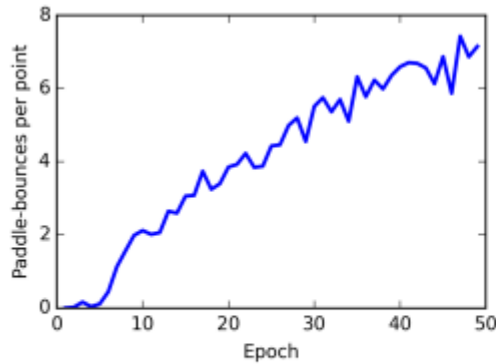
	Left player scores	Right player scores
Left player reward	$\rho$	-1
Right player reward	-1	$\rho$

- Three measures
  - Average paddle-bounces per point
  - Average wall-bounces per paddle-bounce
  - Average serving time per point


# Multi-Agent RL

## □ IQL


### ➤ Result—Fully Competitive



(a) Paddle-bounces per point



(b) Wall-bounces per paddle-bounce




(c) Serving time per point




# Multi-Agent RL

## □ IQL


### ➤ Result—Fully Cooperative





(a) Paddle-bounces per point



(b) Wall-bounces per paddle-bounce



(c) Serving time per point





# Multi-Agent RL

## □ IQL

### ➤ Result—Progression from Competition to Cooperation

Agent	Average paddle-bounces per point	Average wall-bounces per paddle-bounce	Average serving time per point
Competitive $\rho = 1$	$7.15 \pm 1.01$	$0.87 \pm 0.08$	$113.87 \pm 40.30$
Transition $\rho = 0.75$	$7.58 \pm 0.71$	$0.83 \pm 0.06$	$129.03 \pm 38.81$
Transition $\rho = 0.5$	$6.93 \pm 0.49$	$0.64 \pm 0.03$	$147.69 \pm 41.02$
Transition $\rho = 0.25$	$4.49 \pm 0.43$	$1.11 \pm 0.07$	$275.90 \pm 38.69$
Transition $\rho = 0$	$4.31 \pm 0.25$	$0.78 \pm 0.05$	$407.64 \pm 100.79$
Transition $\rho = -0.25$	$5.21 \pm 0.36$	$0.60 \pm 0.05$	$449.18 \pm 99.53$
Transition $\rho = -0.5$	$6.20 \pm 0.20$	$0.38 \pm 0.04$	$433.39 \pm 98.77$
Transition $\rho = -0.75$	$409.50 \pm 535.24$	$0.02 \pm 0.01$	$591.62 \pm 302.15$
Cooperative $\rho = -1$	$654.66 \pm 542.67$	$0.01 \pm 0.00$	$393.34 \pm 138.63$


## □ IL with parameter sharing

- Dec-POMDPs.
- Based on DQN, TRPO, DDPG, A3C
- Three training schemes:
  - **Centralized training and execution**: a centralized policy maps the joint observation of all agents to a joint action
  - **Concurrent training with decentralized execution**: each agent learns its own individual policy.
  - **Parameter sharing during training with decentralized execution**: it allows the policy to be trained with the experiences of all agents simultaneously.


# Multi-Agent RL

## □ IL with parameter sharing


- Four multi-agent benchmark tasks
  - Discrete: Pursuit
  - Continuous: Waterworld, Multi-Walker, Multi-Ant




(a) Pursuit



(b) Waterworld



(c) Multi-Walker




(d) Multi-Ant

# Multi-Agent RL

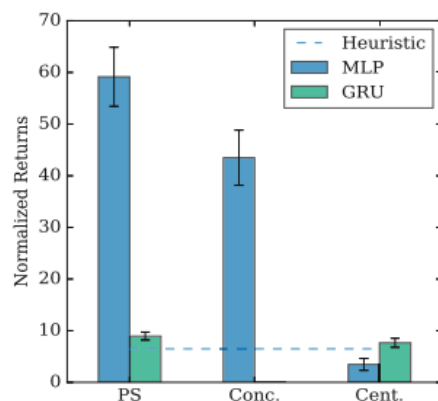
## □ IL with parameter sharing

### ➤ Result


	TRPO	DDPG/DQN	A3C
Feature Net	100-50-25	400-300	128
Recurrent	GRU-32	NA	LSTM-128
Activation	tanh	ReLU	tanh




(a) Pursuit



(b) Waterworld



(c) Multi-Walker



(d) Multi-Ant

Gupta, Jayesh K., Maxim Egorov, and Mykel Kochenderfer. "Cooperative multi-agent control using deep reinforcement learning." *International conference on autonomous agents and multiagent systems*. Springer, Cham, 2017.

# Multi-Agent RL

## □ IL with parameter sharing

### ➤ Result

Task	PS-DQN/DDPG	PS-A3C	PS-TRPO
Pursuit	$10.1 \pm 6.3$	$25.5 \pm 5.4$	$17.4 \pm 4.9$
Waterworld	NA	$10.1 \pm 5.7$	$49.1 \pm 5.7$
Multiwalker	$-8.3 \pm 3.2$	$12.4 \pm 6.1$	$58.0 \pm 4.2$
Multi-ant	$307.2 \pm 13.8$	$483.4 \pm 3.4$	$488.1 \pm 1.3$

- ❑ Independent Learning

- ❑ IQL

- ❑ IL with parameter sharing

- ❑ Learning cooperation

- ❑ MADDPG

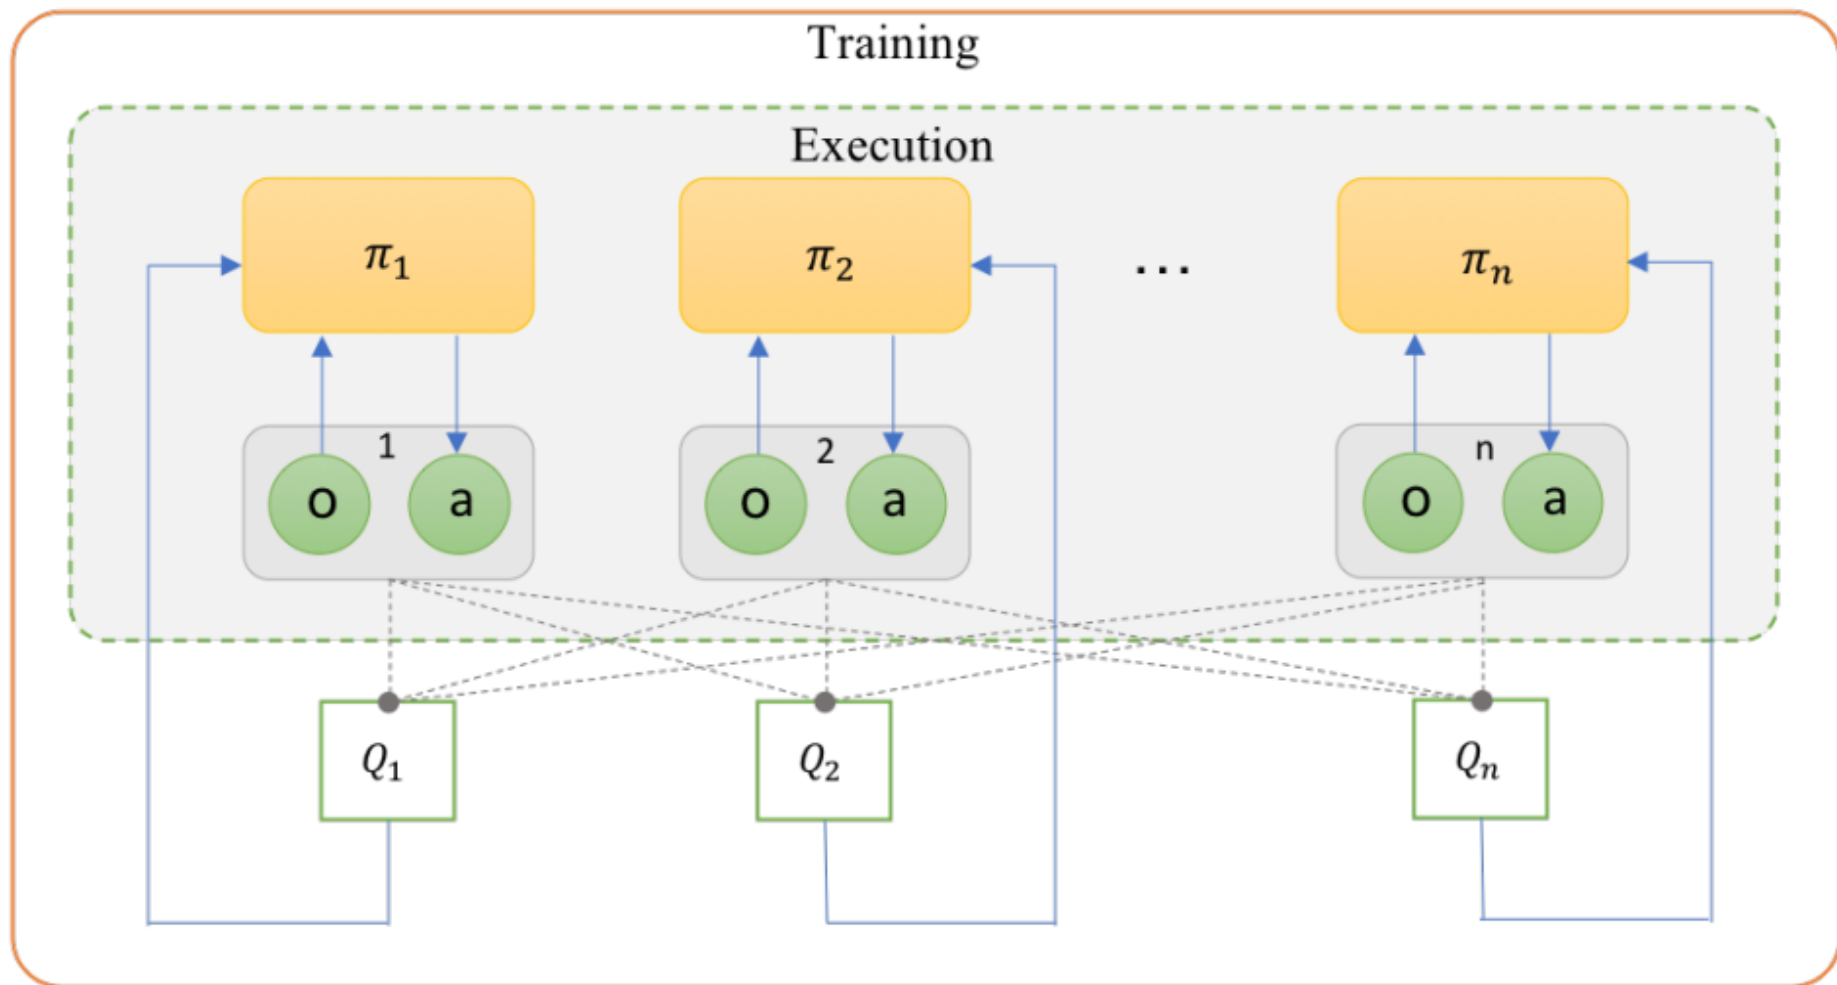
- ❑ COMA

- ❑ VDN

- ❑ QMIX

# Multi-Agent RL

## Centralized Train and Decentralized execution(CTDE)



## □ MADDPG

### ➤ Challenge

- Non-stationarity of the environment
- Policy gradient suffers from a variance that increases as the number of agents grows

### ➤ Main idea

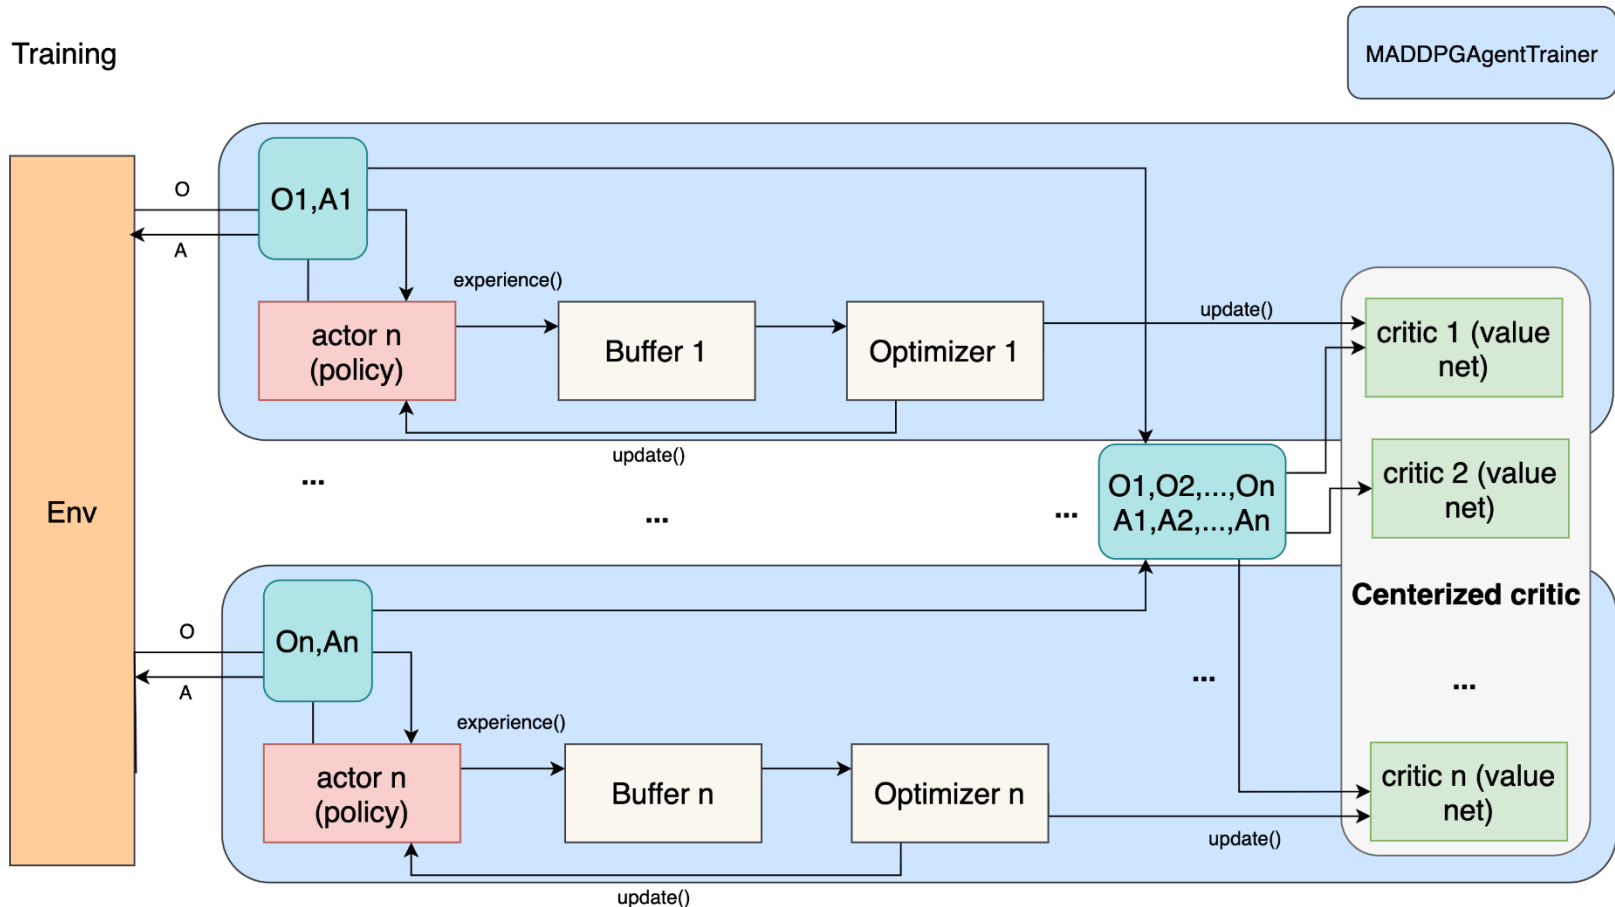
- Leads to learned policies that only use local information at execution time and allow the policies to use extra information to ease training
- Does not assume a differentiable model of the environment dynamics or any particular structure on the communication method between agents
- Is applicable not only to cooperative interaction but to competitive or mixed interaction involving both physical and communicative behavior



# Multi-Agent RL

## □ MADDPG

- Simple extension of actor-critic policy gradient methods where critic is augmented with extra information about the policies of other agents, while the actor only has access to local information



# Multi-Agent RL

## □ MADDPG

- The gradient of the expected return for agent  $i$ :

$$\nabla_{\theta_i} J(\theta_i) = \mathbb{E}_{s \sim p^\mu, a_i \sim \pi_i} \left[ \nabla_{\theta_i} \log \pi_i(a_i | o_i) Q_i^\pi(\mathbf{x}, a_1, \dots, a_N) \right], \mathbf{x} = (o_1, \dots, o_n)$$

- Deterministic policies:

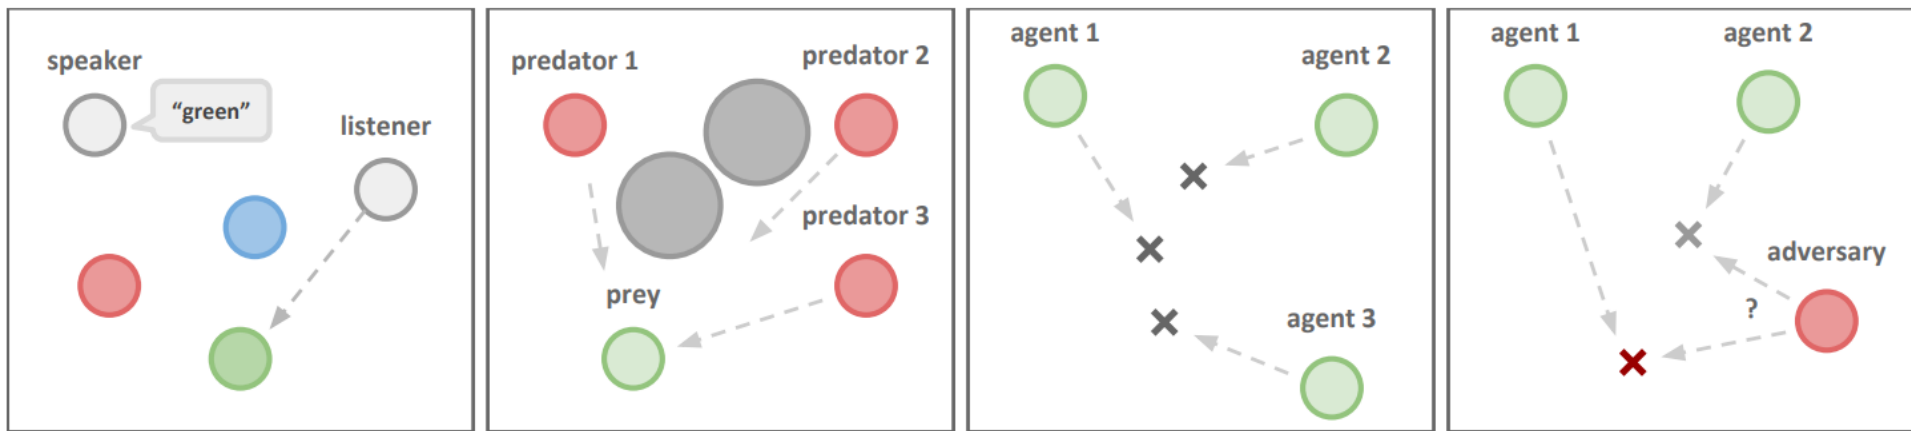
$$\nabla_{\theta_i} J(\mu_i) = \mathbb{E}_{\mathbf{x}, a \sim \mathcal{D}} \left[ \nabla_{\theta_i} \mu_i(a_i | o_i) \nabla_{a_i} Q_i^\mu(\mathbf{x}, a_1, \dots, a_N) \Big|_{a_i = \mu_i(o_i)} \right]$$

$$\mathcal{L}(\theta_i) = \mathbb{E}_{\mathbf{x}, a, r, \mathbf{x}'} \left[ (Q_i^\mu(\mathbf{x}, a_1, \dots, a_N) - y)^2 \right], y = r_i + \gamma Q_i^{\mu'}(\mathbf{x}', a'_1, \dots, a'_N) | a'_j = \mu'_j(o_j)$$

# Multi-Agent RL

## □ MADDPG




- Multiagent-particle-environment
  - Cooperative Communication
  - Predator-Prey
  - Cooperative Navigation
  - Physical Deception
- The environments are publicly available:  
<https://github.com/openai/multiagent-particle-envs>



# Multi-Agent RL

## □ MADDPG

### ➤ Performance



Lowe, Ryan, et al. "Multi-agent actor-critic for mixed cooperative-competitive environments." *Advances in neural information processing systems* 30 (2017).

## □ COMA

### ➤ Challenge

- Credit assignment: in cooperative settings, joint actions typically generate only global rewards, making it difficult for each agent to deduce its own contribution to the team's success.

### ➤ Main idea

- Use CTDE as same as MADDPG
- **Counterfactual: Use a counterfactual baseline that marginalizes out a single agent's action**
  - **Using the centralized critic to compute an agent-specific advantage function** that compares the estimated return for the current joint action to a counterfactual baseline that marginalizes out a single agent's action, while keeping the other agent's actions fixed
- Uses a critic representation that allows the counterfactual baseline to be computed efficiently.

# Multi-Agent RL

## □ COMA

➤ Advantage function:

$$A^a(s, \mathbf{u}) = Q(s, \mathbf{u}) - \sum \pi^a(u'^a | \tau^a) Q(s, (\mathbf{u}^{-a}, u'^a))$$

➤ COMA gradient:

$$g = \mathbb{E}_\pi \left[ \sum_a \nabla_\theta \log \pi^a(u^a | \tau^a) A^a(s, \mathbf{u}) \right], A^a(s, \mathbf{u}) = Q(s, \mathbf{u}) - b(s, \mathbf{u}^{-a})$$

$$g_b = -\mathbb{E}_\pi \left[ \sum_a \nabla_\theta \log \pi^a(u^a | \tau^a) b(s, \mathbf{u}^{-a}) \right]$$

$$g_b = -\sum_s d^\pi(s) \sum_a \sum_{\mathbf{u}^{-a}} \pi(\mathbf{u}^{-a} | \tau^{-a}) \sum_{u^a} \pi^a(u^a | \tau^a) \nabla_\theta \log \pi^a(u^a | \tau^a) b(s, \mathbf{u}^{-a})$$

$$= -\sum_s d^\pi(s) \sum_a \sum_{\mathbf{u}^{-a}} \pi(\mathbf{u}^{-a} | \tau^{-a}) \sum_{u^a} \nabla_\theta \pi^a(u^a | \tau^a) b(s, \mathbf{u}^{-a})$$

$$= -\sum_s d^\pi(s) \sum_a \sum_{\mathbf{u}^{-a}} \pi(\mathbf{u}^{-a} | \tau^{-a}) b(s, \mathbf{u}^{-a}) \nabla_\theta 1$$

$$= 0$$


# Multi-Agent RL

## COMA


➤ COMA gradient:

$$g = \mathbb{E}_{\pi} \left[ \sum_a \nabla_{\theta} \log \pi^a(u^a | \tau^a) A^a(s, \mathbf{u}) \right], A^a(s, \mathbf{u}) = Q(s, \mathbf{u}) - b(s, \mathbf{u}^{-a})$$


$$= \mathbb{E}_{\pi} \left[ \sum_a \nabla_{\theta} \log \pi^a(u^a | \tau^a) Q(s, \mathbf{u}) \right] = \mathbb{E}_{\pi} \left[ \nabla_{\theta} \log \prod_a \pi^a(u^a | \tau^a) Q(s, \mathbf{u}) \right]$$



(a)



(b)



(c)

# Multi-Agent RL

## □ COMA

---

### Algorithm 1 Counterfactual Multi-Agent (COMA) Policy Gradients

---

```

Initialise  $\theta_1^c, \hat{\theta}_1^c, \theta^\pi$ 
for each training episode  $e$  do
  Empty buffer
  for  $e_c = 1$  to  $\frac{\text{BatchSize}}{n}$  do
     $s_1$  = initial state,  $t = 0$ ,  $h_0^a = \mathbf{0}$  for each agent  $a$ 
    while  $s_t \neq \text{terminal}$  and  $t < T$  do
       $t = t + 1$ 
      for each agent  $a$  do
         $h_t^a = \text{Actor}(o_t^a, h_{t-1}^a, u_{t-1}^a, a, u; \theta_i)$ 
        Sample  $u_t^a$  from  $\pi(h_t^a, \epsilon(e))$ 
        Get reward  $r_t$  and next state  $s_{t+1}$ 
      Add episode to buffer
    Collate episodes in buffer into single batch
    for  $t = 1$  to  $T$  do // from now processing all agents in parallel via single batch
      Batch unroll RNN using states, actions and rewards
      Calculate TD( $\lambda$ ) targets  $y_t^a$  using  $\hat{\theta}_i^c$ 
    for  $t = T$  down to 1 do
       $\Delta Q_t^a = y_t^a - Q(s_t^a, \mathbf{u})$ 
       $\Delta \theta^c = \nabla_{\theta^c} (\Delta Q_t^a)^2$  // calculate critic gradient
       $\theta_{i+1}^c = \theta_i^c - \alpha \Delta \theta^c$  // update critic weights
      Every C steps reset  $\hat{\theta}_i^c = \theta_i^c$ 
    for  $t = T$  down to 1 do
       $A^a(s_t^a, \mathbf{u}) = Q(s_t^a, \mathbf{u}) - \sum_u Q(s_t^a, u, \mathbf{u}^{-a}) \pi(u|h_t^a)$  // calculate COMA
       $\Delta \theta^\pi = \Delta \theta^\pi + \nabla_{\theta^\pi} \log \pi(u|h_t^a) A^a(s_t^a, \mathbf{u})$  // accumulate actor gradients
     $\theta_{i+1}^\pi = \theta_i^\pi + \alpha \Delta \theta^\pi$  // update actor weights

```


---




# Multi-Agent RL

## COMA

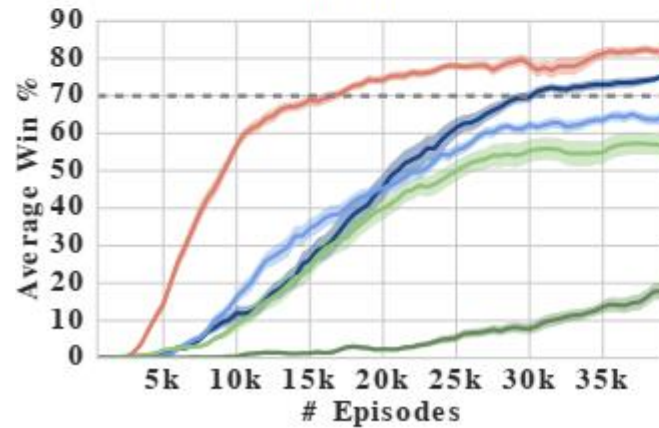
### ➤ Performance




(a) 3m



(b) 5m



(c) 5w



(d) 2d\_3z

## □ VDN

### ➤ Challenge

- Lazy agent: one agent learns a useful policy, but a second agent is discouraged from learning because its exploration would hinder the first agent and lead to worse team reward.
- Large-scale multi-agent scenarios.

### ➤ Main idea

- Training individual agents with a novel value decomposition network architecture, which learns to decompose the team value function into agent-wise value functions
- The value decomposition network aims to learn an optimal linear value decomposition from the team reward signal, by back-propagating the total Q gradient through deep neural networks representing the individual component value functions.

## □ VDN

- Joint action-value function can be additively decomposed into value functions across agents

$$Q((h^1, \dots, h^d), (a^1, \dots, a^d)) \approx \sum_{i=1}^d \tilde{Q}_i(h^i, a^i)$$

- Consider the case with two agents and where rewards decompose additively across agent observations,  $r(\mathbf{s}, \mathbf{a}) = r_1(o_t^1, a_t^1) + r_2(o_t^2, a_t^2)$

$$\begin{aligned} Q^\pi(\mathbf{s}, \mathbf{a}) &= \mathbb{E} \left[ \sum_{t=1}^{\infty} \gamma^{t-1} r(\mathbf{s}_t, \mathbf{a}_t) \mid \mathbf{s}_1 = \mathbf{s}, \mathbf{a}_1 = \mathbf{a}; \pi \right] \\ &= \mathbb{E} \left[ \sum_{t=1}^{\infty} \gamma^{t-1} r_1(o_t^1, a_t^1) \mid \mathbf{s}_1 = \mathbf{s}, \mathbf{a}_1 = \mathbf{a}; \pi \right] + \mathbb{E} \left[ \sum_{t=1}^{\infty} \gamma^{t-1} r_2(o_t^2, a_t^2) \mid \mathbf{s}_1 = \mathbf{s}, \mathbf{a}_1 = \mathbf{a}; \pi \right] \\ &=: \bar{Q}_1^\pi(\mathbf{s}, \mathbf{a}) + \bar{Q}_2^\pi(\mathbf{s}, \mathbf{a}) \end{aligned}$$

- When agents store additional information from historical observation

$$Q^\pi(\mathbf{s}, \mathbf{a}) =: \bar{Q}_1^\pi(\mathbf{s}, \mathbf{a}) + \bar{Q}_2^\pi(\mathbf{s}, \mathbf{a}) \approx \tilde{Q}_1^\pi(h^1, a^1) + \tilde{Q}_2^\pi(h^2, a^2)$$

# Multi-Agent RL

## □ VDN

### ➤ Architecture

IL

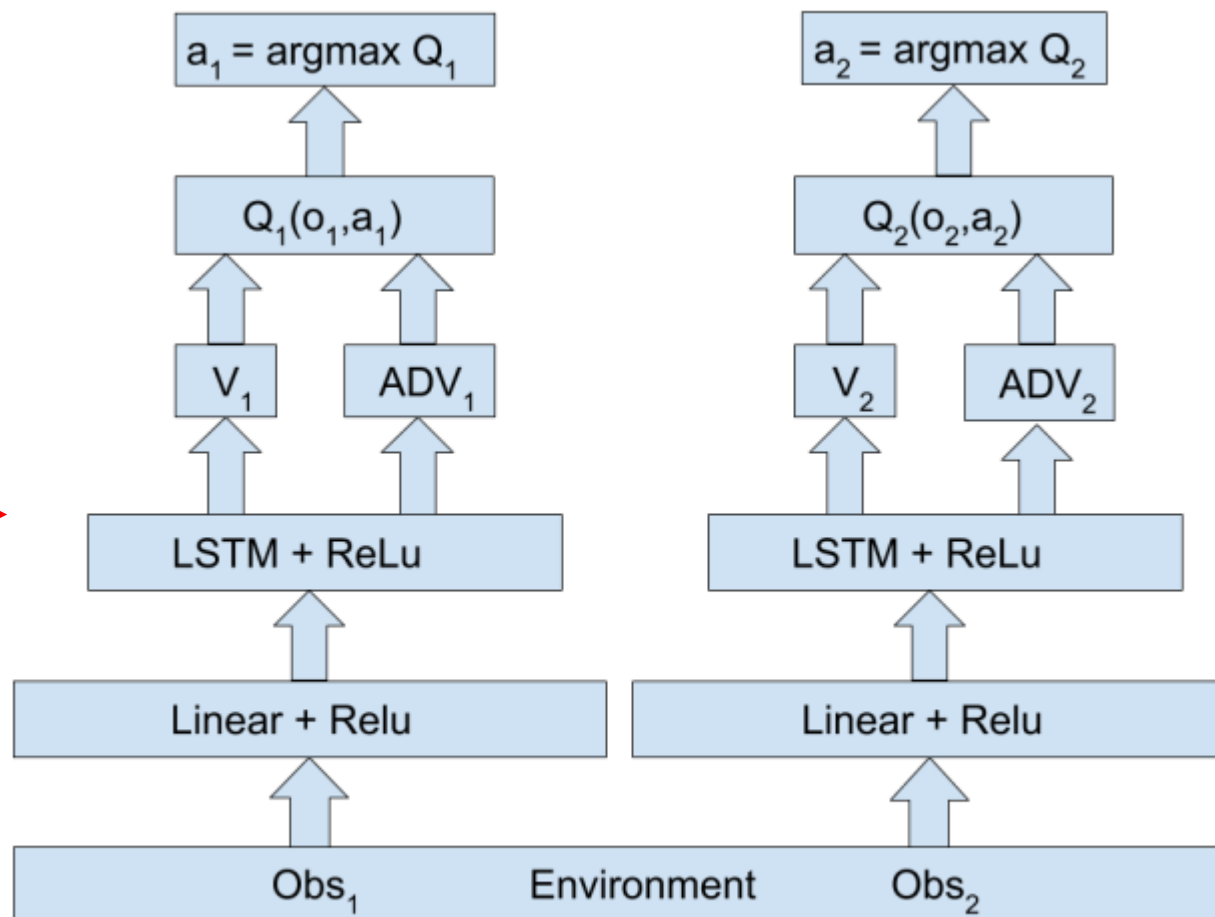


Figure 14: Independent Agents Architecture

# Multi-Agent RL

## □ VDN

### ➤ Architecture

VDN →




Figure 15: Value-Decomposition Individual Architecture

# Multi-Agent RL

## □ VDN

### ➤ Architecture

VDN+Comm →

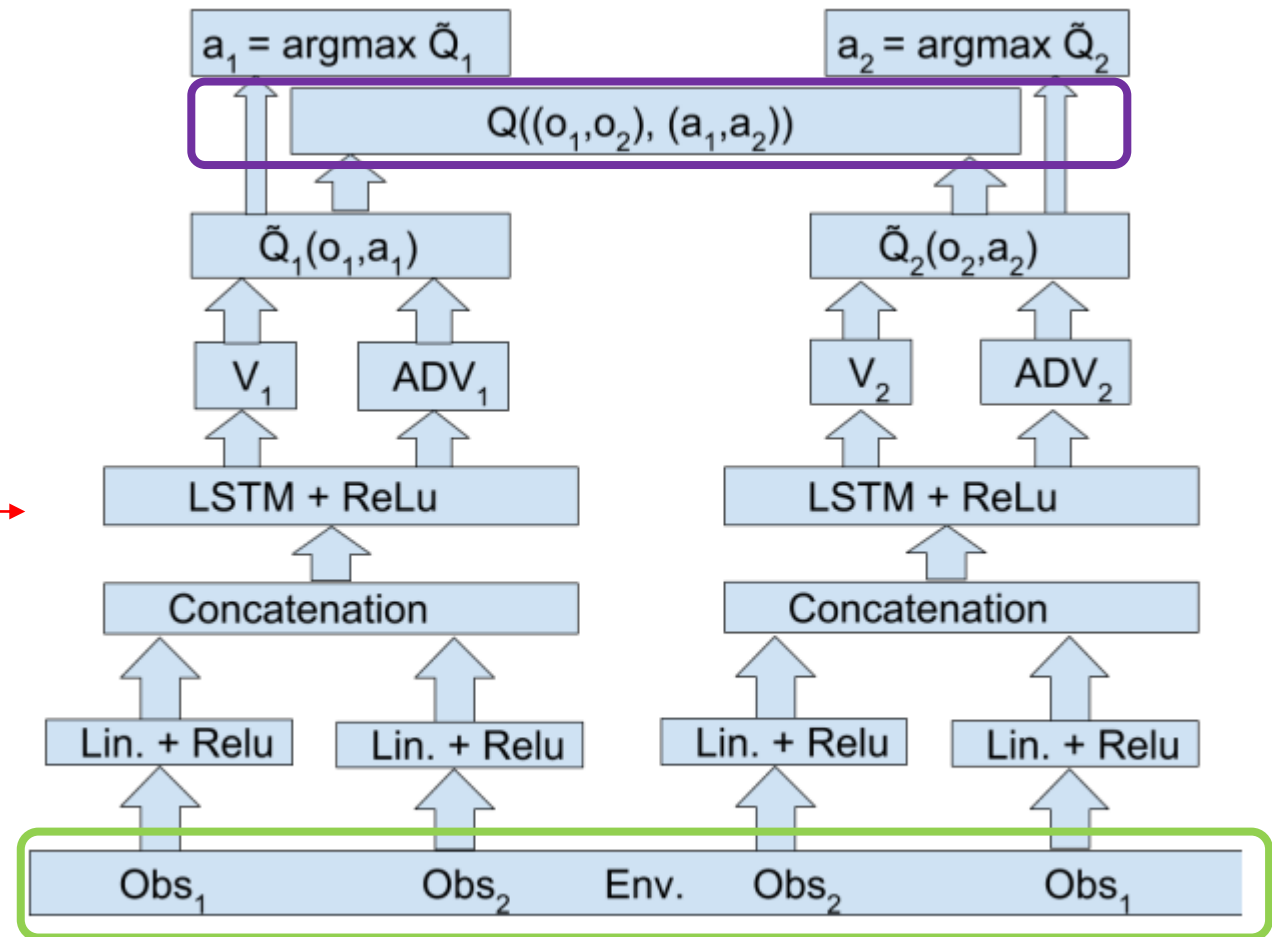


Figure 16: Low-level communication Architecture

# Multi-Agent RL

## □ VDN

### ➤ Architecture

VDN+Comm →




Figure 17: High-level communication Architecture

# Multi-Agent RL

## □ VDN

### ➤ Architecture

Centralized →

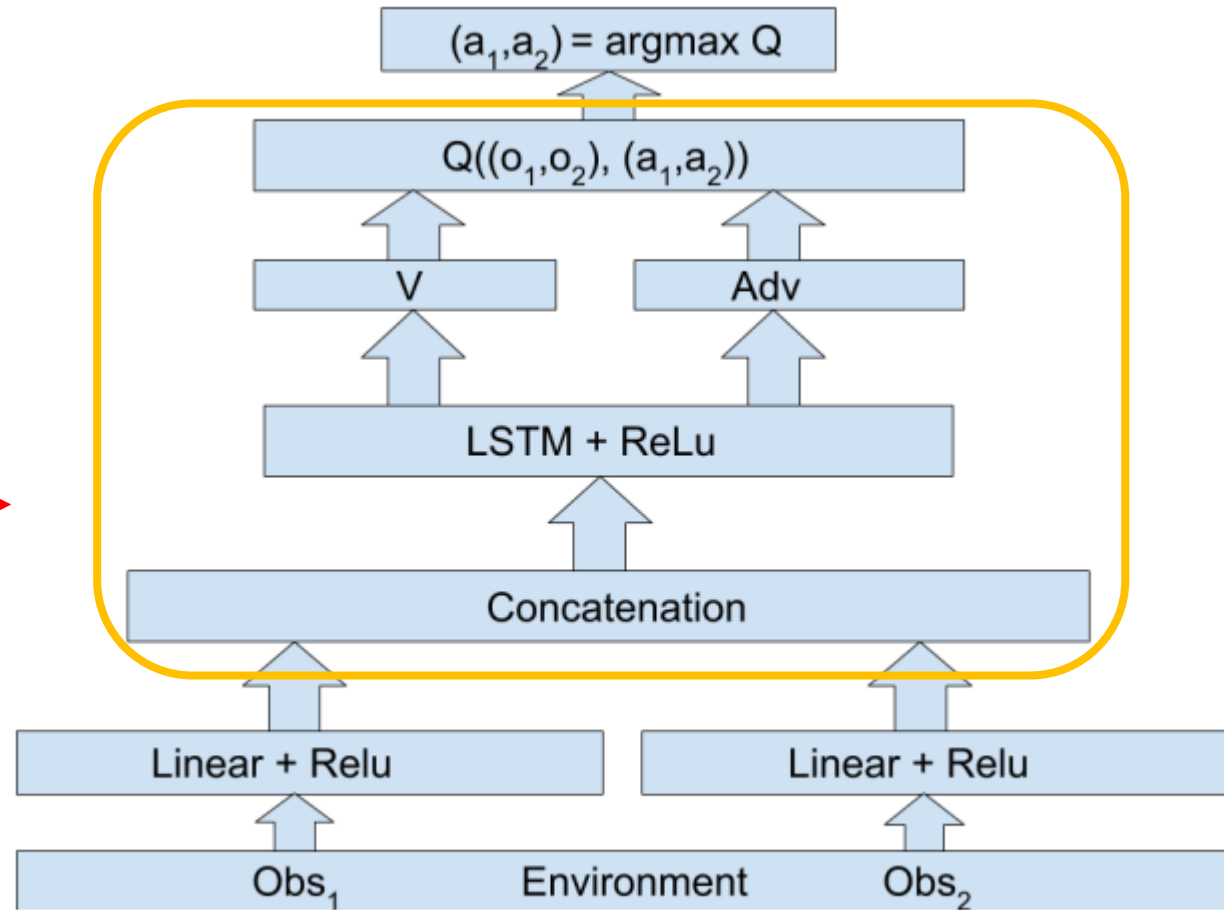


Figure 20: Combinatorially Centralized Architecture



# Multi-Agent RL

## □ VDN

### ➤ Performance

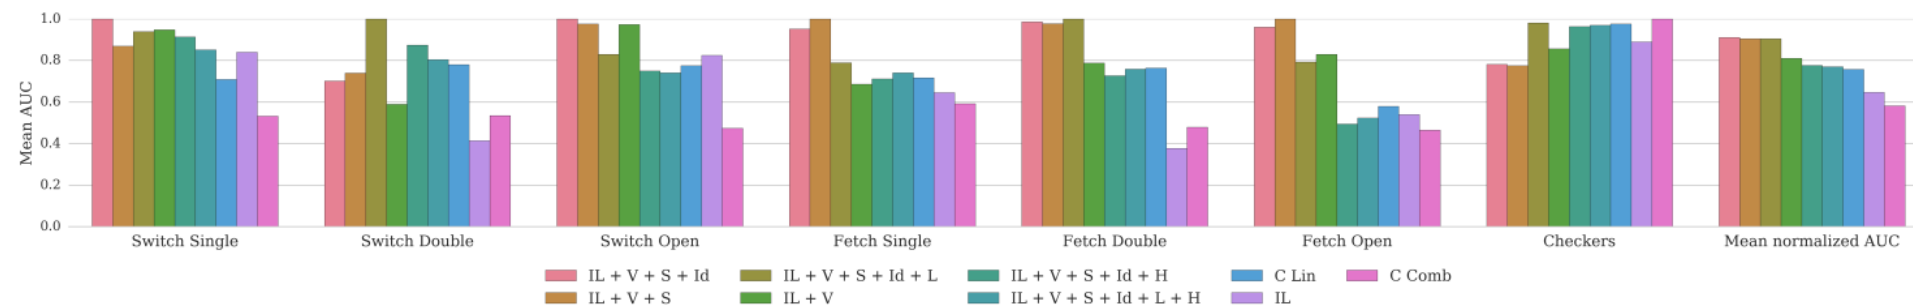
Agent	V.	S.	Id	L.	H.	C.
1						
2	✓					
3	✓	✓				
4	✓	✓	✓			
5	✓	✓	✓	✓		
6	✓	✓	✓		✓	
7	✓	✓	✓	✓	✓	
8	✓					✓
9						✓

Table 1: Agent architectures. V is value decomposition, S means shared weights and an invariant network, Id means role info was provided, L stands for lower-level communication, H for higher-level communication and C for centralization. These architectures were selected to show the advantages of the independent agent with value-decomposition and to study the benefits of additional enhancements added in a logical sequence.


# Multi-Agent RL

## VDN

### ➤ Performance



Architecture	IL + V + S + Id + L	0.96	0.98	0.79	0.82	1.00	0.84	0.97	0.91
	IL + V + S + Id	0.78	1.00	0.85	0.96	0.69	1.00	1.00	0.90
	IL + V + S	0.75	0.90	1.00	1.00	0.63	0.96	0.91	0.88
	IL + V	0.88	0.86	0.76	0.70	0.66	0.94	0.98	0.83
	IL + V + S + Id + L + H	0.95	0.82	0.49	0.78	0.80	0.72	0.85	0.77
	IL + V + S + Id + H	0.95	0.70	0.45	0.72	0.86	0.68	0.93	0.76
	C Lin	0.96	0.86	0.47	0.73	0.68	0.76	0.72	0.74
	IL	0.92	0.38	0.47	0.66	0.44	0.82	0.83	0.64
	C Comb	1.00	0.46	0.38	0.53	0.60	0.45	0.54	0.57
	Task	Checkers	Fetch Double	Fetch Open	Fetch Single	Switch Double	Switch Open	Switch Single	Mean final reward



# Multi-Agent RL

## □ QMIX

### ➤ Main idea

- Employ a network that estimates joint action-values as a complex non-linear combination of per-agent values that condition only on local observations

- Ensure that a global argmax performed on  $Q_{tot}$  yields the same result as a set of individual argmax operations performed on each  $Q_a$

$$\operatorname{argmax}_{\mathbf{u}} Q_{tot}(\boldsymbol{\tau}, \mathbf{u}) = \begin{pmatrix} \operatorname{argmax}_{u^1} Q_1(\tau^1, u^1) \\ \vdots \\ \operatorname{argmax}_{u^n} Q_n(\tau^n, u^n) \end{pmatrix}$$

- Enforce a monotonicity constraint on the relationship between  $Q_{tot}$  and each  $Q_a$

$$\frac{\partial Q_{tot}}{\partial Q_a} \geq 0, \forall a \in A$$

## □ QMIX

### ➤ Main idea

- Enforce a monotonicity constraint on the relationship between  $Q_{tot}$  and each  $Q_a$

$$\frac{\partial Q_{tot}}{\partial Q_a} \geq 0, \forall a \in A$$

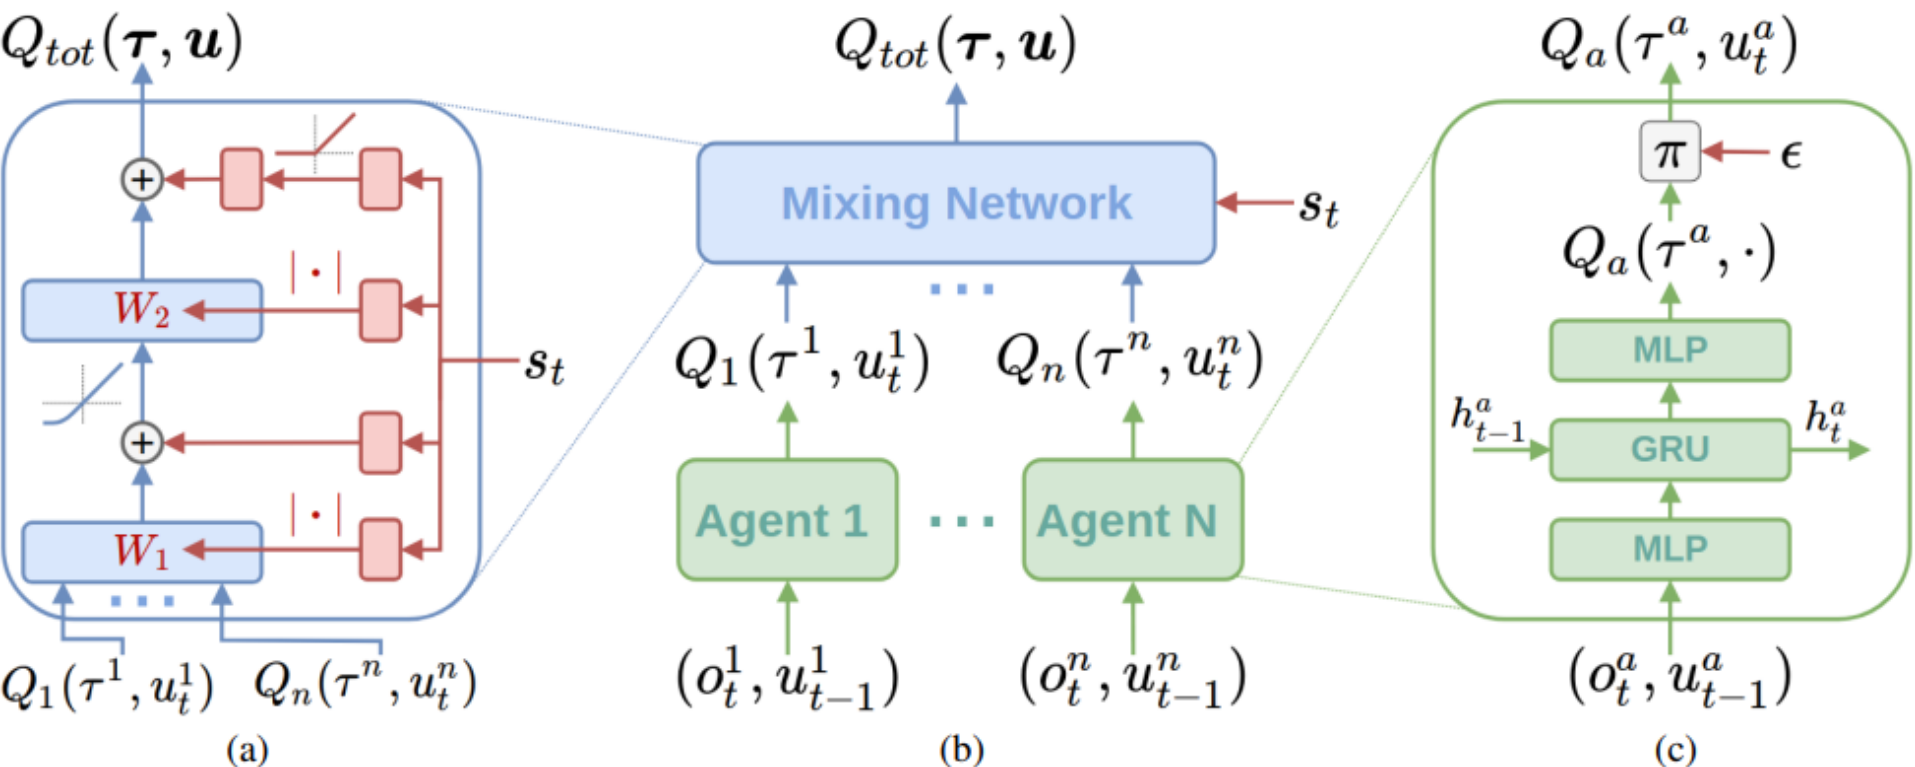
- Represent  $Q_{tot}$  using an architecture consisting of agent networks, a mixing network, and a set of hypernetworks.
- Restrict the mixing network to have positive weights
- The weights of the mixing network are produced by separate hypernetworks.
- Each hypernetwork takes the state  $s$  as input and generates the weights of one layer of the mixing network. Each hypernetwork consists of a single linear layer, followed by an absolute activation function, to ensure that the mixing network weights are non-negative.

# Multi-Agent RL

## □ QMIX

- Architecture
- Loss

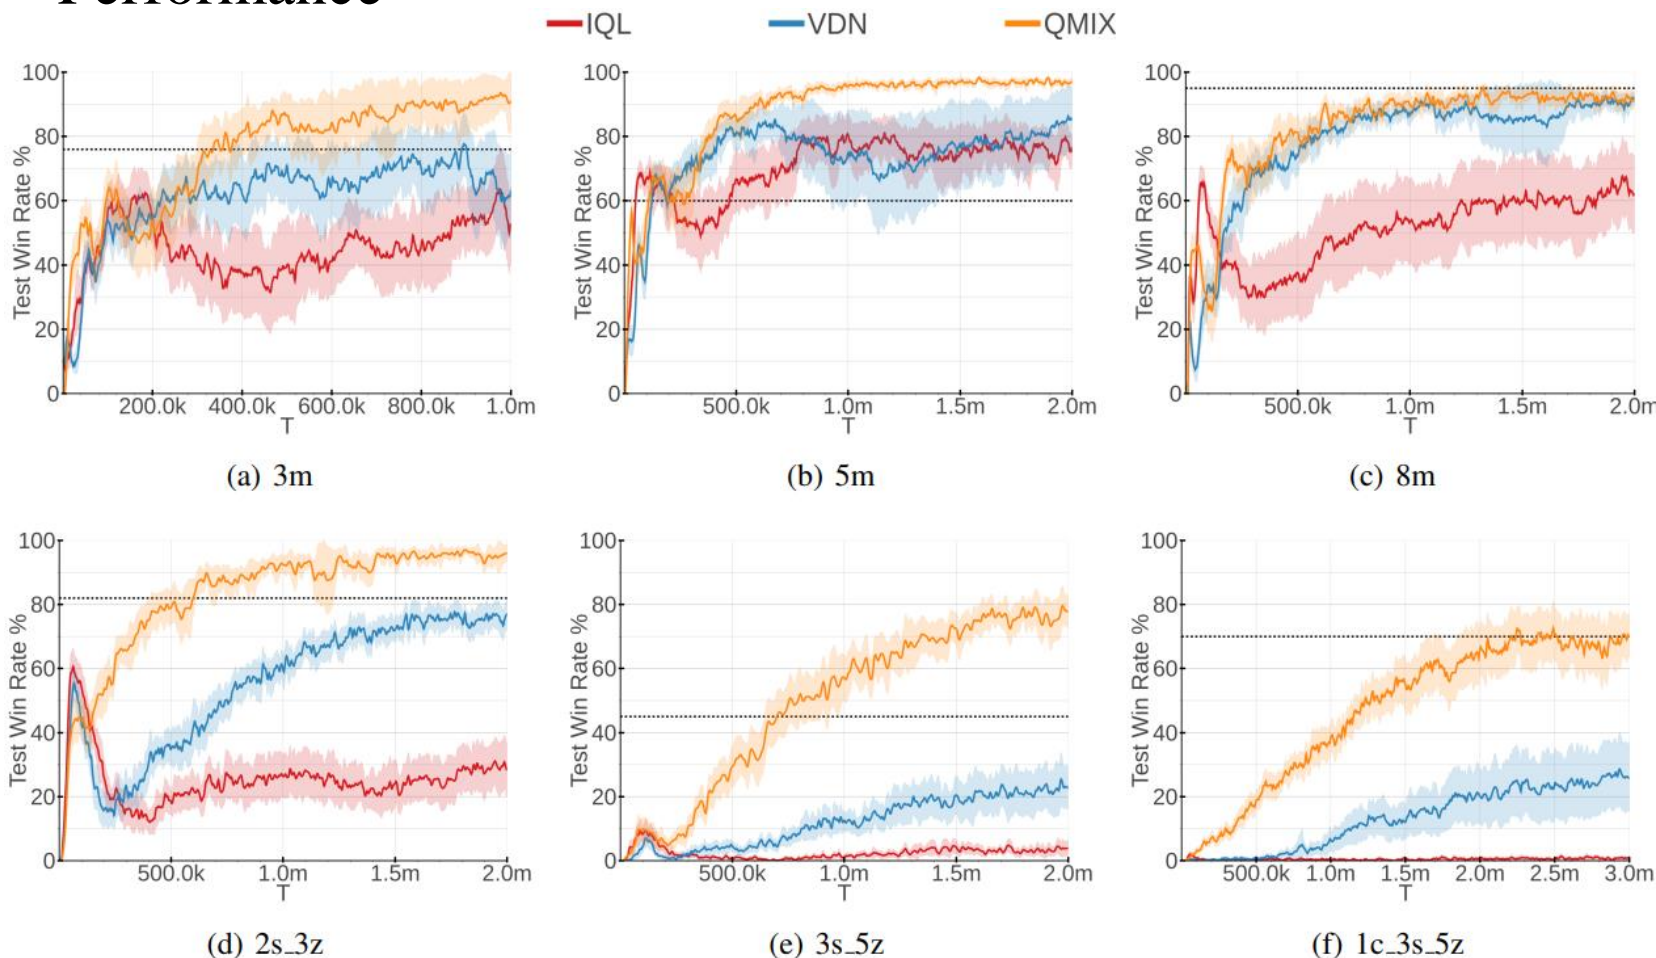
$$\mathcal{L}(\theta) = \sum_{i=1}^b \left[ \left( y_i^{tot} - Q_{tot}(\tau, \mathbf{u}, s; \theta) \right)^2 \right]$$



# Multi-Agent RL

## □ QMIX

### ➤ Performance



Rashid, Tabish, et al. "Qmix: Monotonic value function factorisation for deep multi-agent reinforcement learning." *International Conference on Machine Learning*. PMLR, 2018.



# Multi-Agent RL

## □ QMIX

### ➤ Performance

- QMIX-NS: the weights and biases of the mixing network are learned in the standard way, without conditioning on the state
- QMIX-LIN: remove the hidden layer of the mixing network
- VDN-S: add a state-dependent term to the sum of the agent's Q-values

